Application Security Best Practices and Scaling Applications on Google Cloud

Student: Islam Aip

Course: Cloud application development

Date: 24.11.2024

Almaty, 2024

Introduction

Application Security Best Practices

1. Google Cloud Project

```
aipislam03@cloudshell:~$ gcloud projects create f1g2h3
Create in progress for [https://cloudresourcemanager.googleapis.com/v1/projects/f1g2h3].
Waiting for [operations/cp.6911677119749184565] to finish...done.
Enabling service [cloudapis.googleapis.com] on project [f1g2h3]...
Operation "operations/acat.p2-573076915860-69b01007-8206-4988-9618-781b049be47a" finished successfully.
aipislam03@cloudshell:~$
```

Let’s create a new GCP project and enable all necessary APIs.

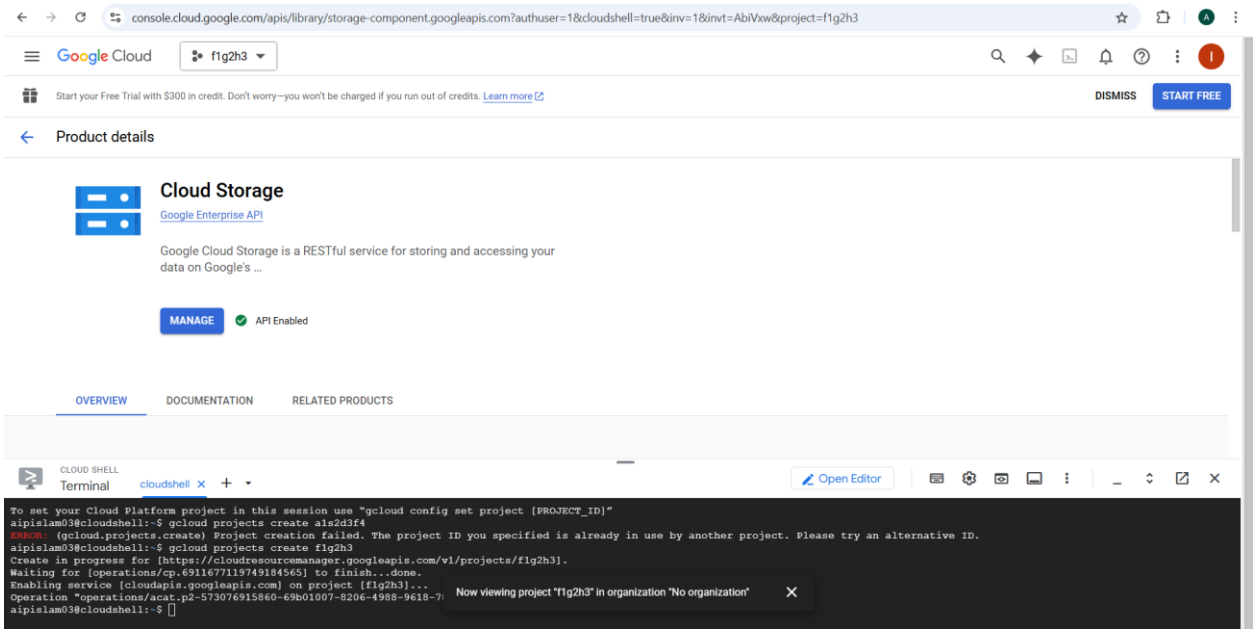

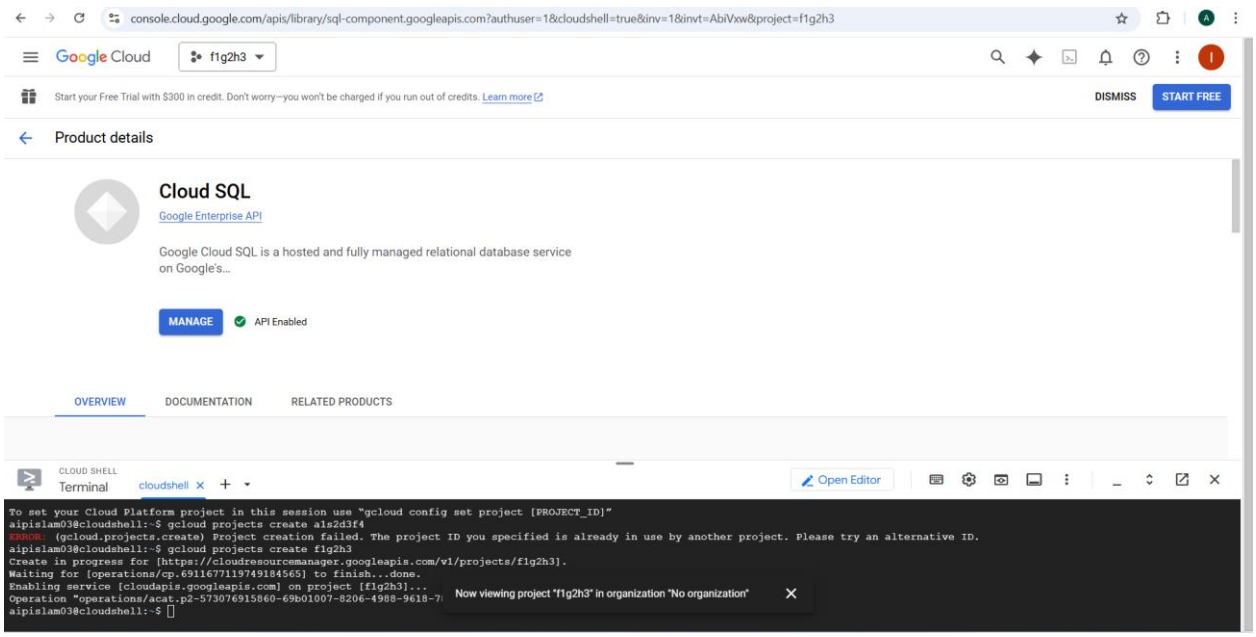

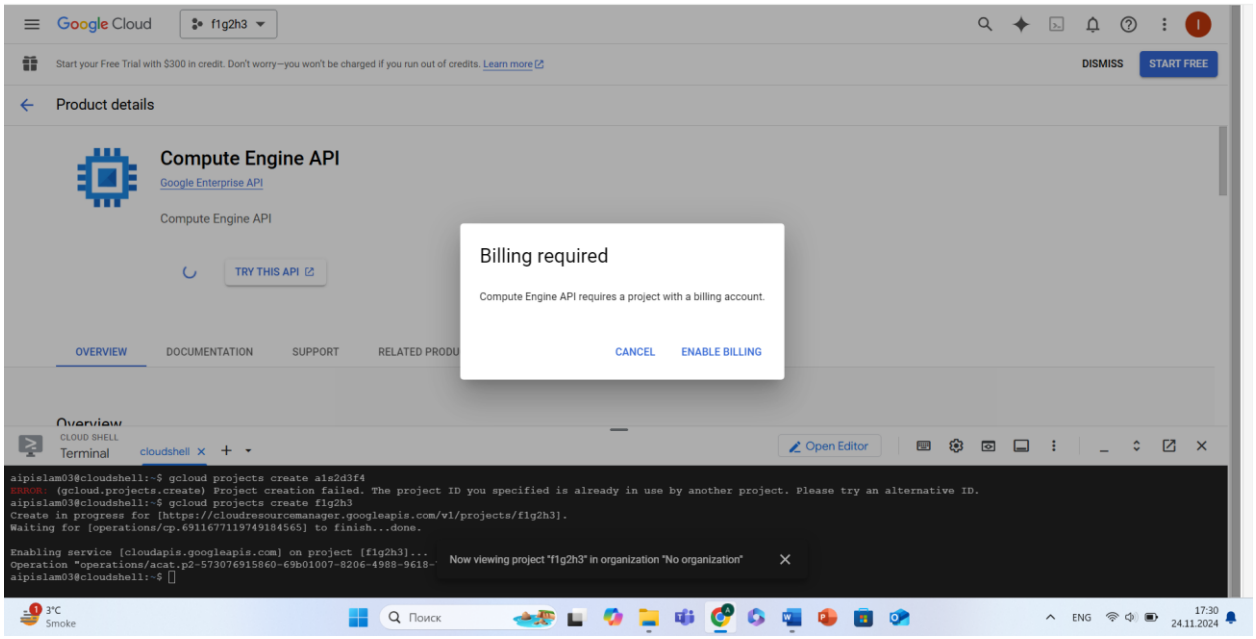

For some APIs billing should be enabled.

2. Identity and Access Management

Now let's create a service account and grant access to project.

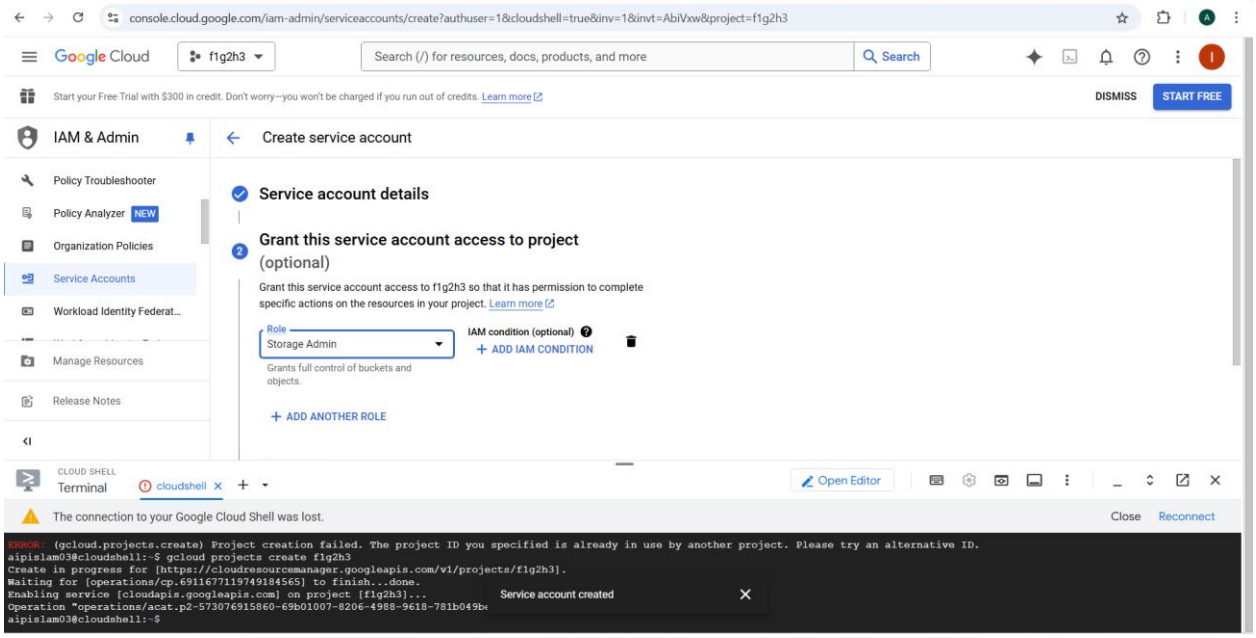

Now we can assign conditions based on attributes such as time, day of week and so on.

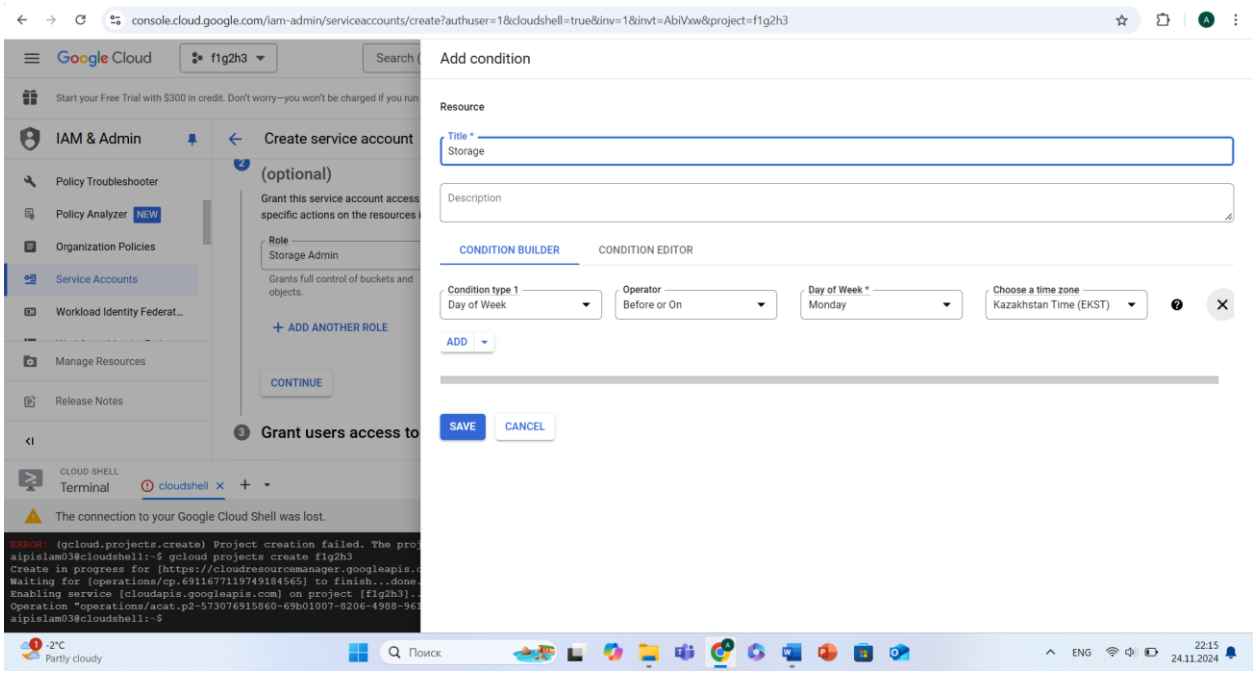

3. Data Protection

In order set up encryption for data at rest we can use Google Cloud KMS.

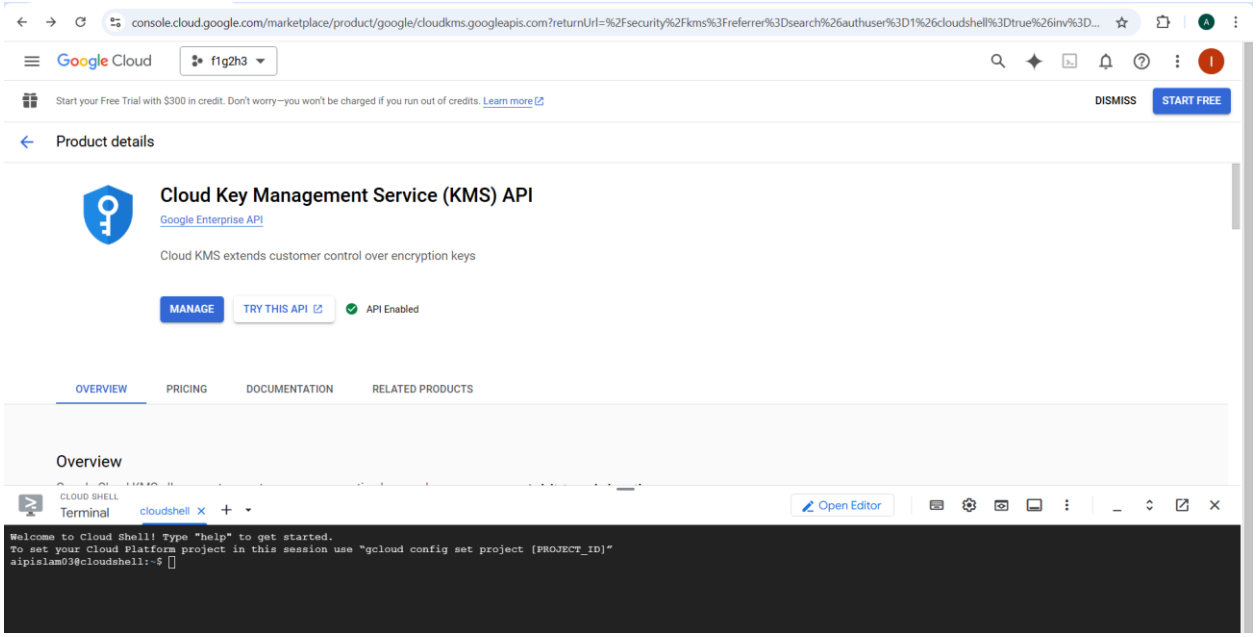

After enabling Google Cloud KMS API, we can create credentials and choose our service account that we created before.

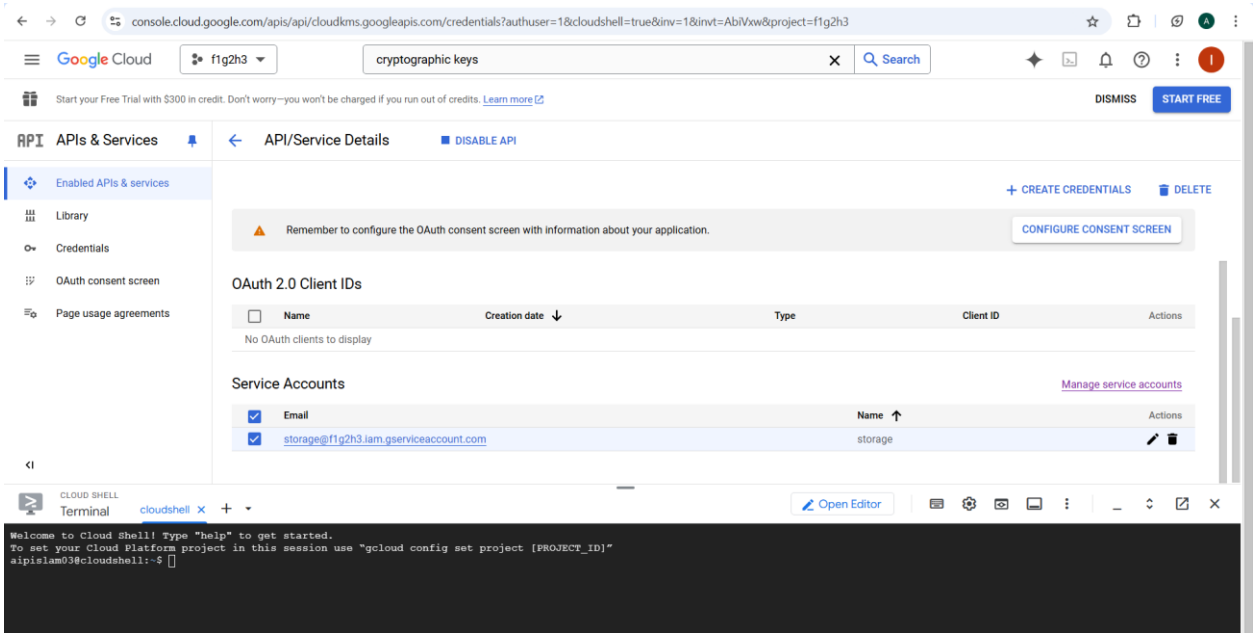

Now, in order to set up a load balancer, we need billing account.

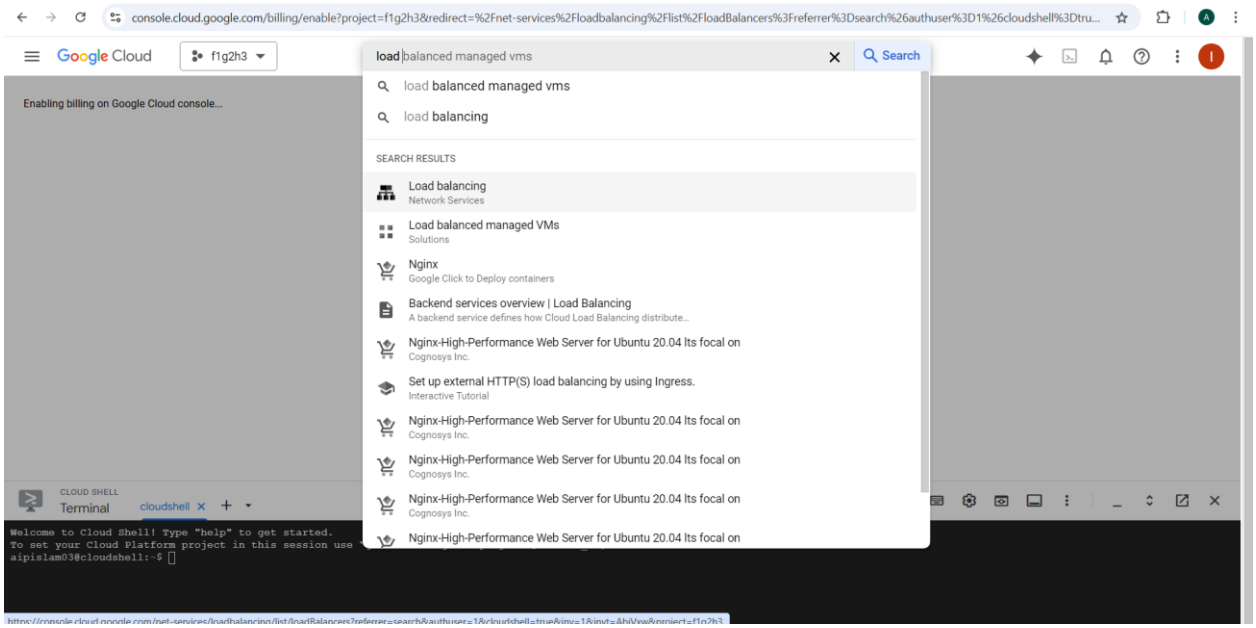

After enabling billing we should be able to configure load balance and upload an SSL certificate or use Google managed SSL.

4. Application Security Testing

Security testing is crucial part of the application development. In Google Cloud we can integrate a security scanning tools such as Snyk, or we can use security command center to check threats, vulnerabilities and risks.

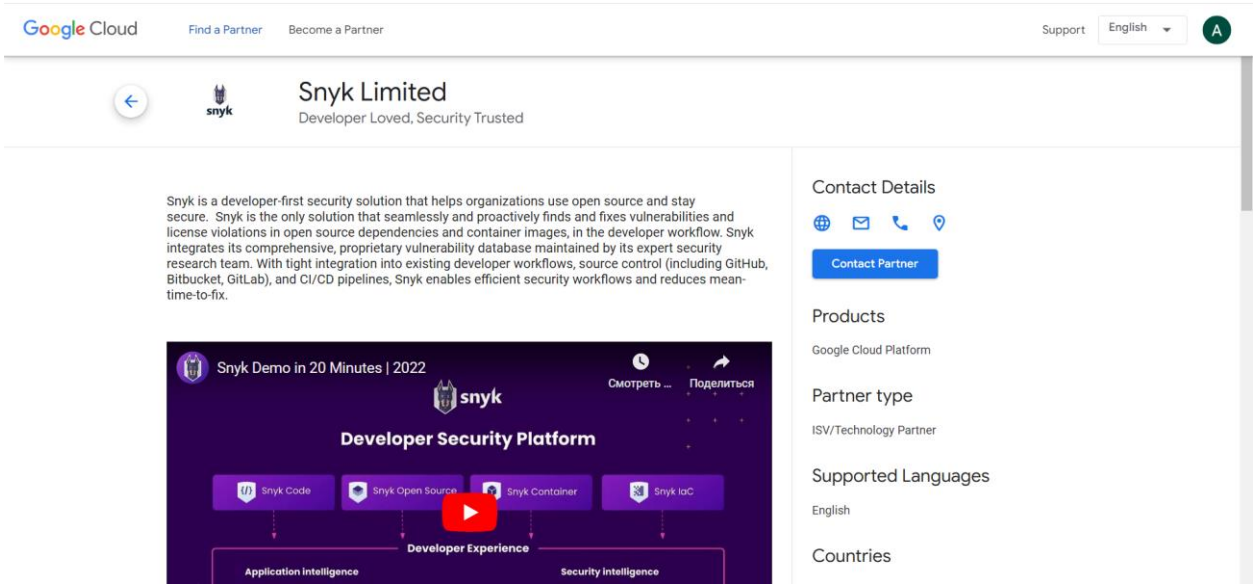

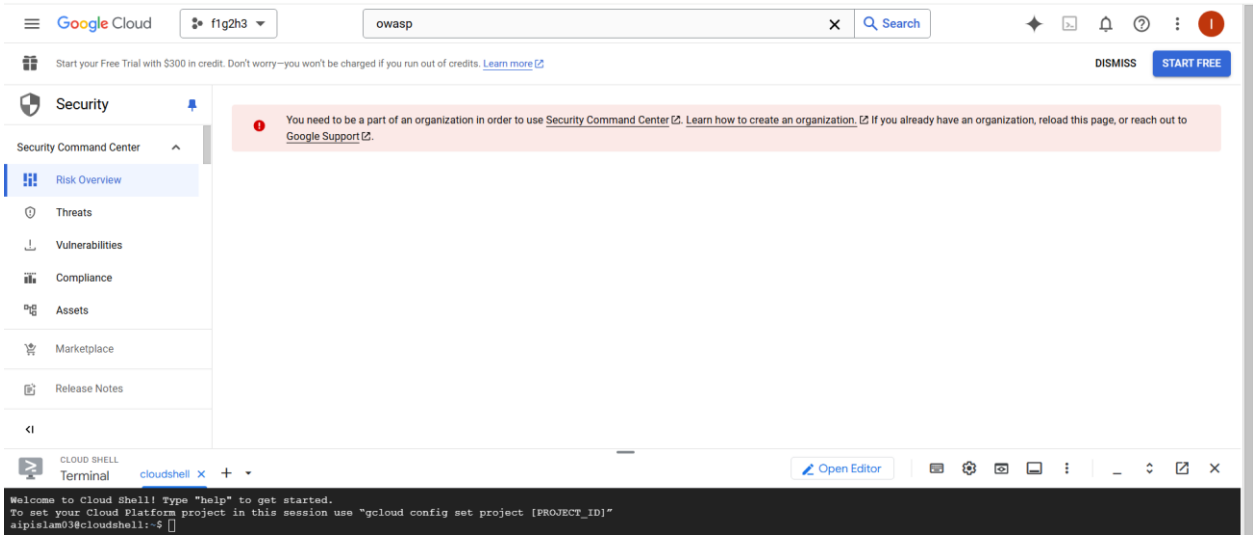

5. Monitoring and Logging

In Google Cloud we can use Logs Explorer for logging.

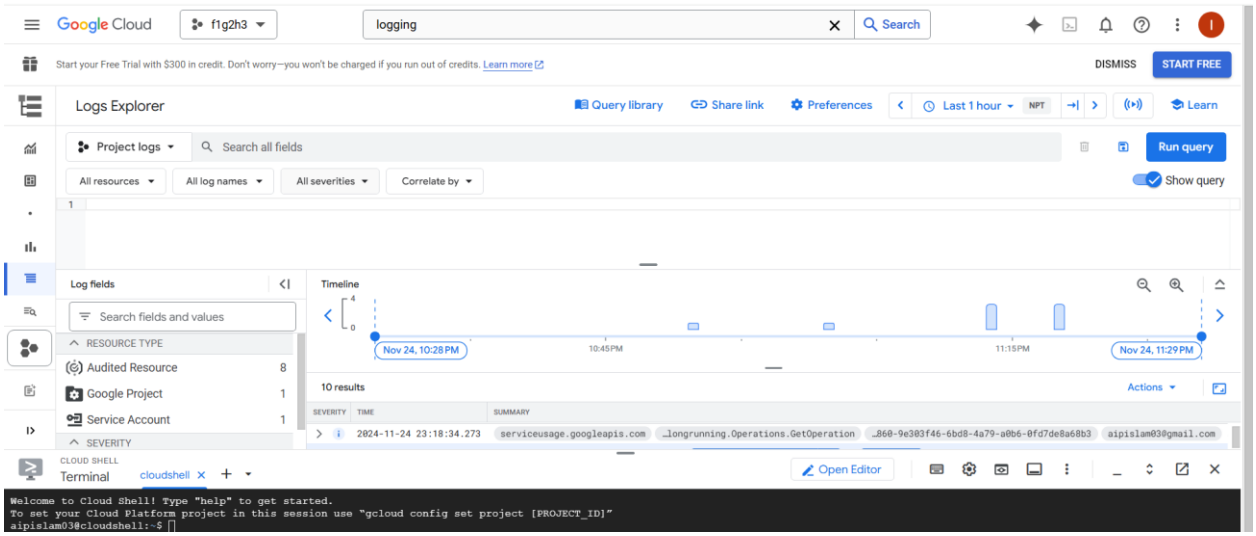

Now let’s enable audit logs. Admin Activity logs are typically enabled by default, for data access logs let’s enable read and write for Cloud Storage.

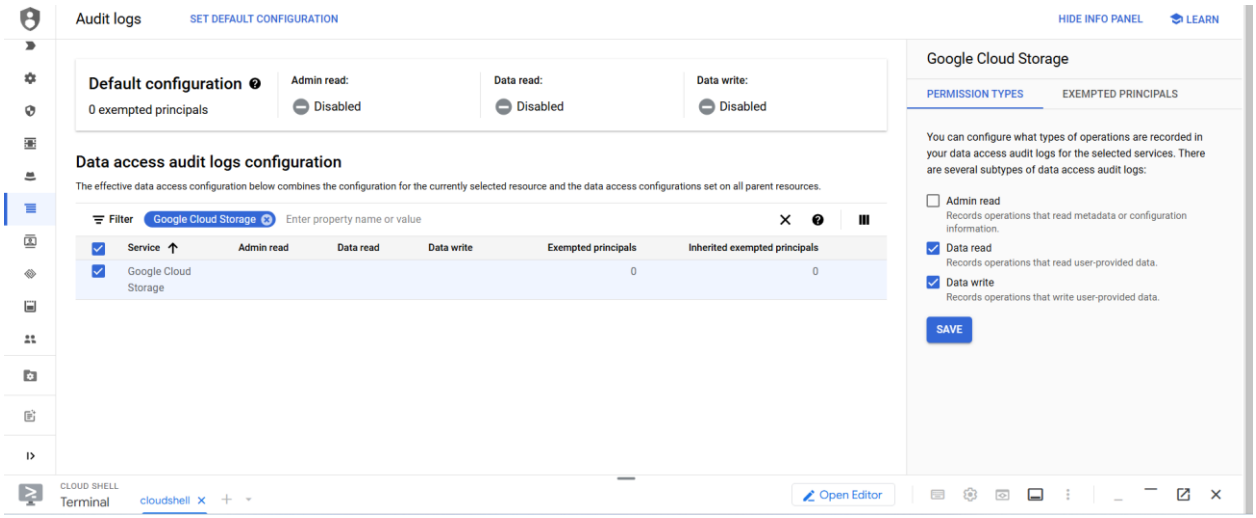

Now, to monitor suspicious activities or errors we can create a new alert policy in Monitoring > Alerting.

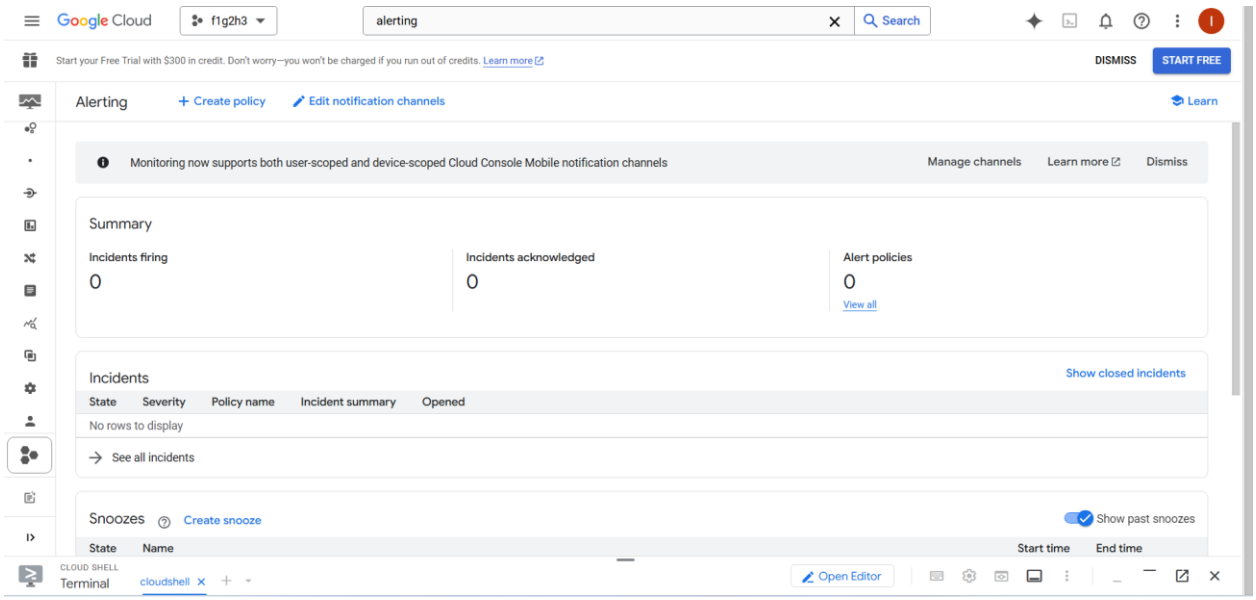

Here we can configure alert trigger for number of log entries.

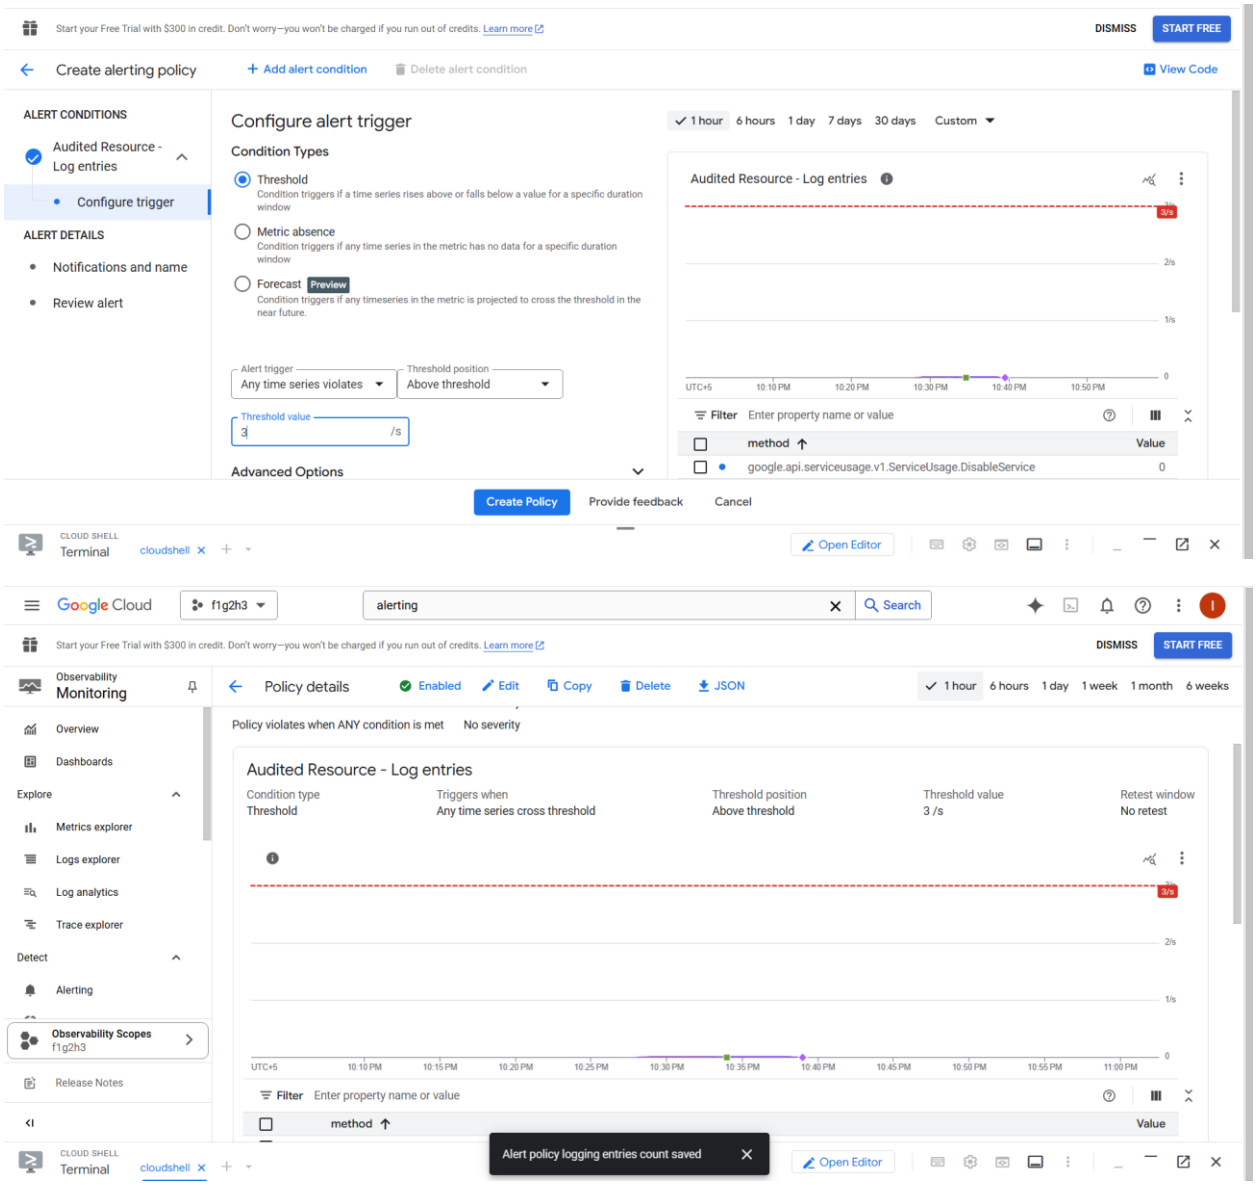

6. Incident Response

An Incident Response Plan ensures that team of developers can effectively handle security breaches. IRPs steps include:

1. Identify – we define how to detect and classify an incident. We can use tools like Cloud monitoring, Cloud Security Command Center, or logs from Cloud Audit Logs.

2. Contain – we define steps to isolate the system to prevent further damage. We can use tool like IAM permissions.

3. Eradicate – we remove the root cause of the incident, such as deleting malicious code or fixing vulnerability. We can patch software, update IAM roles etc.

4. Recover – we restore system to normal operation and ensure that vulnerability is fixed.

Scaling Applications on Google Cloud

1. Application Design

Let’s create a simple to-do app.

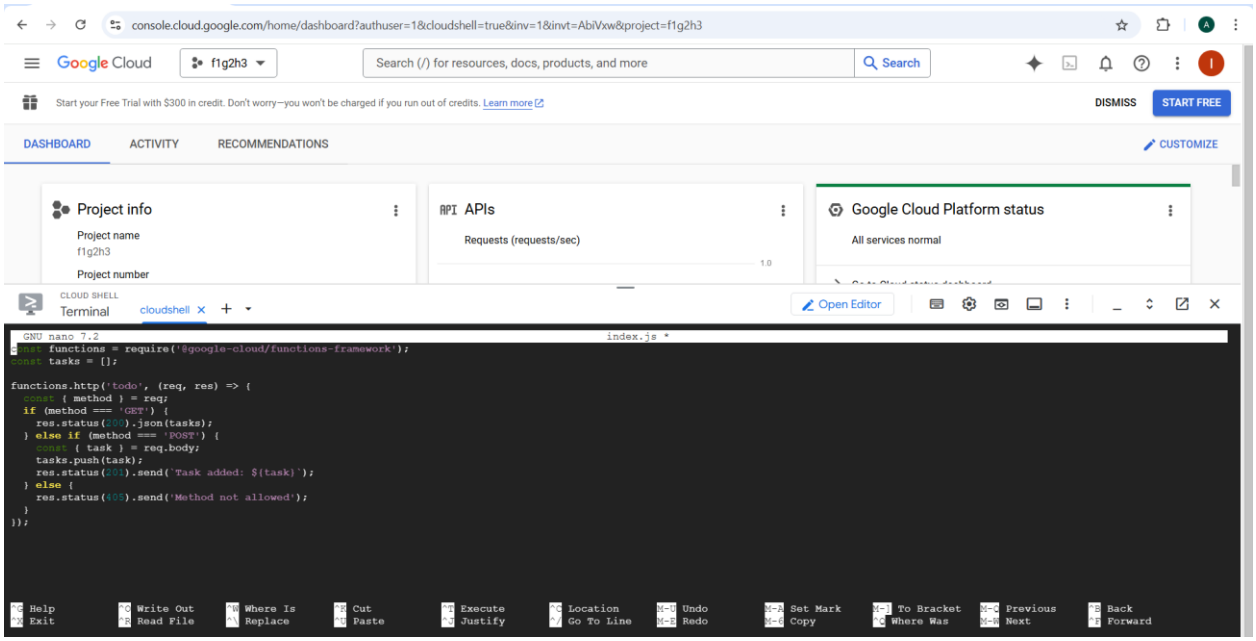

We can use serverless computing or GKE for containerized applications. Let’s add Dockerfile for our app and build the Docker image:

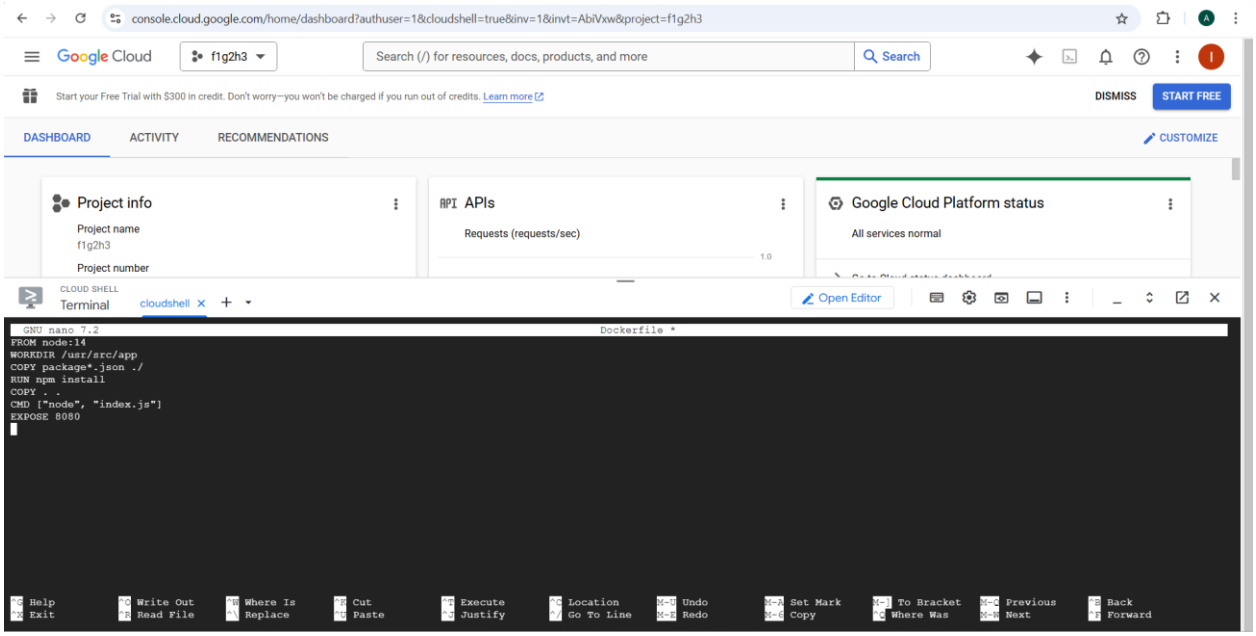

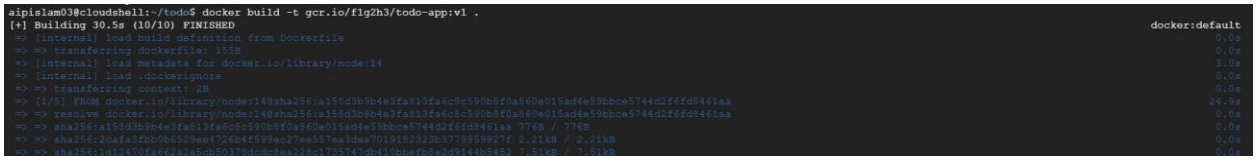

Now we can push the image to the Google Container Registry, but we need billing account to be enabled:

```
aipislam03@cloudshell:~/todo$ docker push gcr.io/f1g2h3/todo-app:v1
The push refers to repository [gcr.io/f1g2h3/todo-app]
e38e5ebcbb96: Preparing
666c60fedfca: Preparing
5f70bf18a086: Preparing
3baf2d7099d9: Preparing
0d5f5a015e5d: Preparing
3c777d951de2: Waiting
f8a91dd5fc84: Waiting
cb81227abde5: Waiting
e01a454893a9: Waiting
c45660adde37: Waiting
fe0fb3ab4a0f: Waiting
f1186e5061f2: Waiting
b2dba7477754: Waiting
denied: Artifact Registry API has not been used in project 573076915860 before or it is disabled. Enable it by visiting https://console.developers.google.com/apis/api/artifactregistry.googleapis.com/overview?project=573076915860 then retry. If you enabled this API recently, wait a few minutes for the action to propagate to our systems and retry.
aipislam03@cloudshell:~/todo$
```

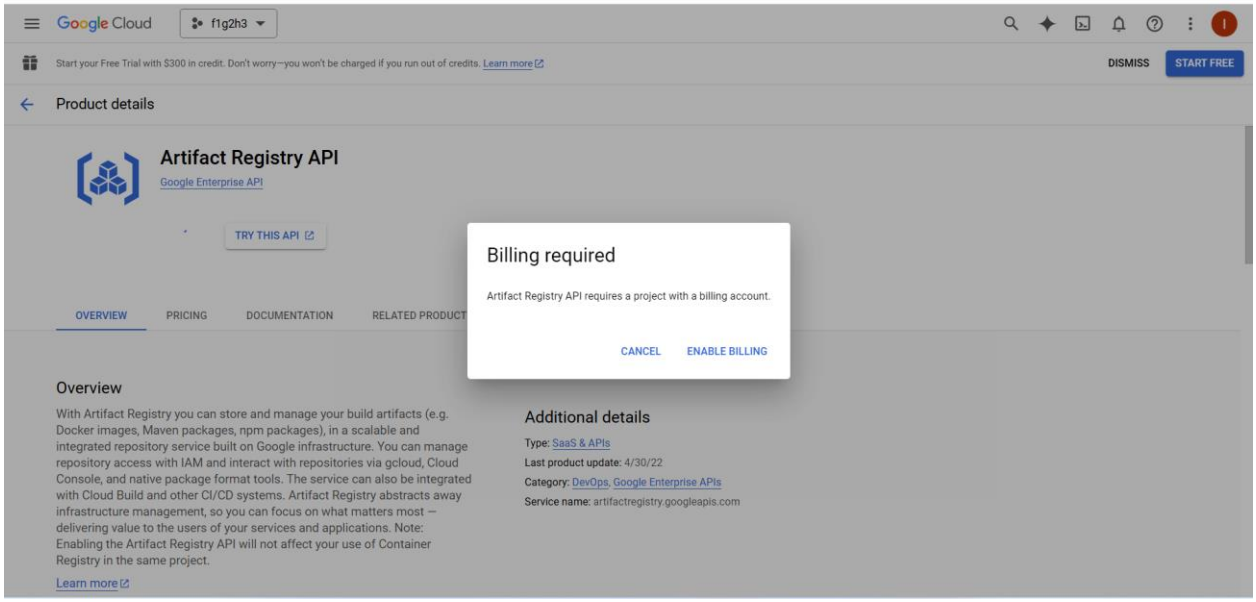

After pushing the image we can deploy our application on GKE.

We create a GKE cluster, deploy container to the cluster and expose our app.

```
aipislam03@cloudshell:~/todo$ gcloud container clusters create todo-cluster \
    --num-nodes=3 \
    --zone=[ZONE]
```

```
aipislam03@cloudshell:~/todo$ kubectl create deployment todo-app --image=gcr.io/f1g2h3/todo-app:v1
```

```
aipislam03@cloudshell:~/todo$ kubectl expose deployment todo-app --type=LoadBalancer --port 80 --target-port 8080
```

2. Horizontal vs. Vertical Scaling

Horizontal scaling involves adding more instances or nodes to distribute the workload. This approach is preferable over vertical scaling in some scenarios:

1) A web application or API must remain available even if one server fails.
2) Application experiences spikes in traffic.
3) Application designed to use multiple smaller nodes instead of a single powerful one.

In GKE we can enable horizontal auto pod scaling:

```
aipislam03@cloudshell:~/todo$ kubectl autoscale deployment todo-app --cpu-percent=80 --min=1 --max=5
```

For Vertical Scaling in Compute Engine we can navigate to our VM instances and upgrade the machine type(billing account is required):

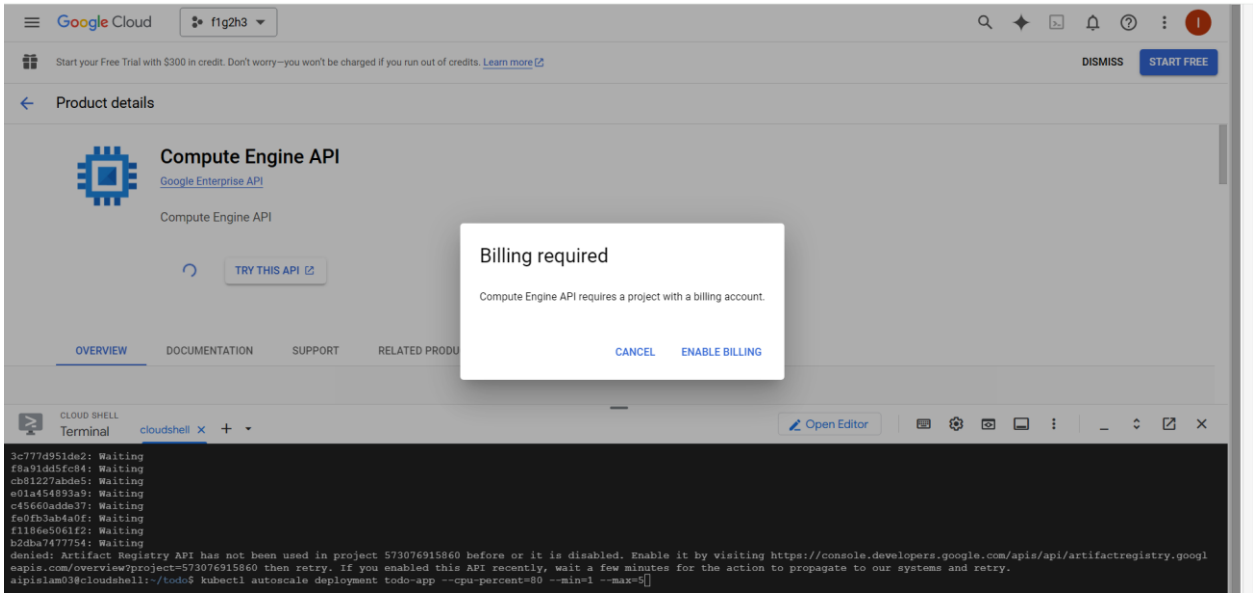

3. Load Balancing

In order to set up a Load Balancer for our project we can navigate to Network Services > Load balancing

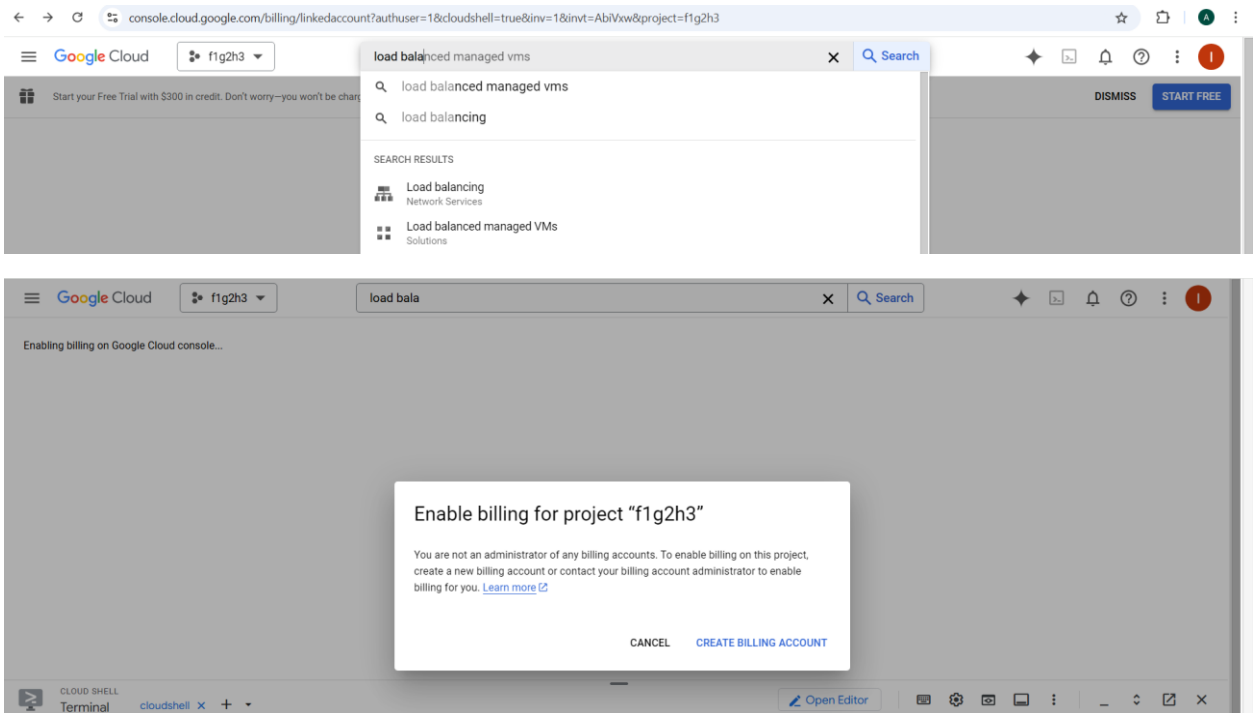

Billing is required

Here we create Load Balancer and we can choose HTTPS Load balancer. After creating, we can attach backend services(GKE pods, Compute Engine instances).

4. Auto-Scaling

For Compute Engine we can create an instance group with auto-scaling ebnabled:

```
aipislam03@cloudshell:~/todo$ gcloud compute instance-groups managed create todo-group \
    --base-instance-name=todo-instance \
    --template=[INSTANCE_TEMPLATE] \
    --size=1 \
    --zone=[ZONE]
gcloud compute instance-groups managed set-autoscaling todo-group \
    --max-num-replicas=10 \
    --min-num-replicas=1 \
    --target-cpu-utilization=0.6 \
    --zone=[ZONE]
```

For GKE:

```
aipislam03@cloudshell:~/todo$ kubectl autoscale deployment todo-app --cpu-percent=80 --min=1 --max=10
```

5. Monitoring Performance

1) Track Metrics

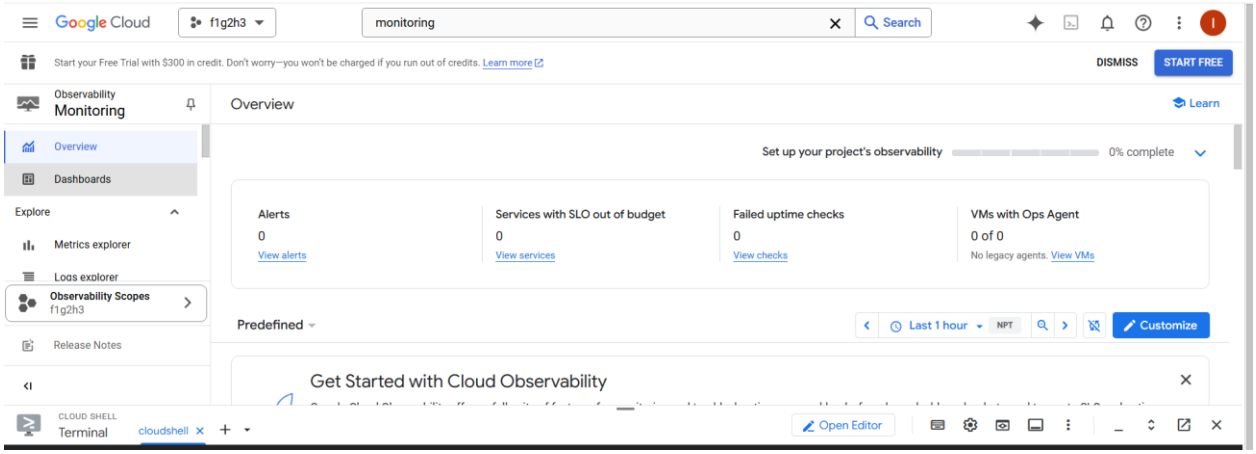

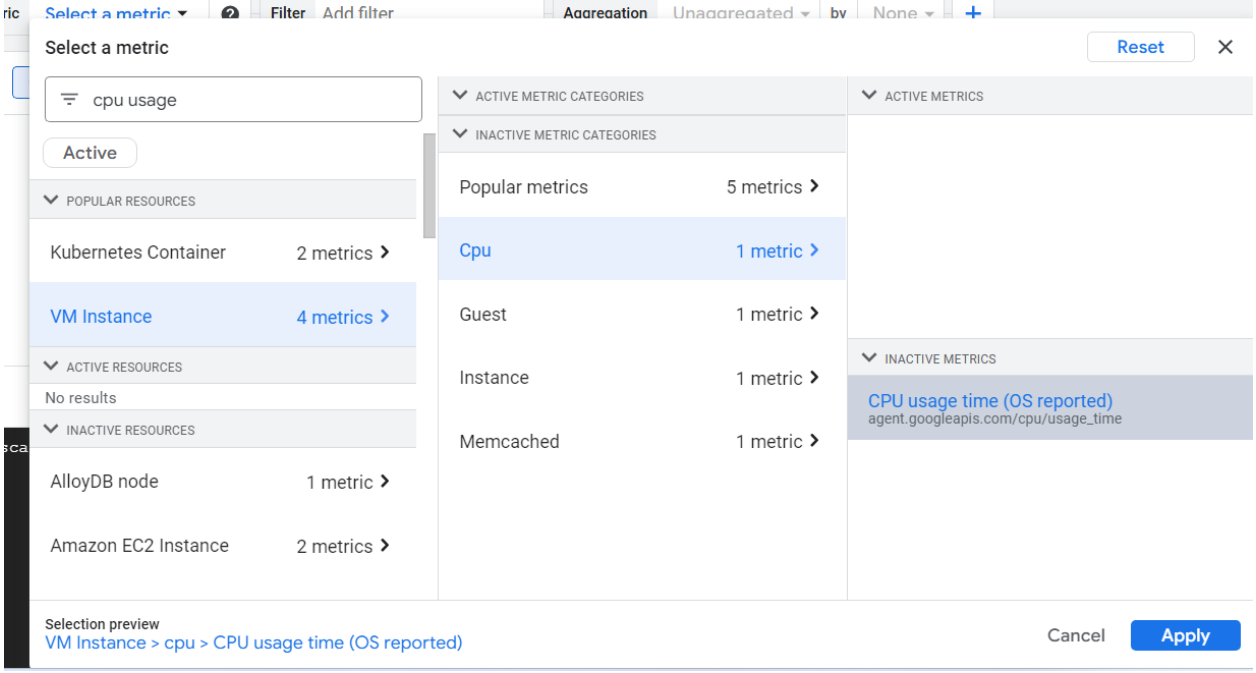

2) Dashboards

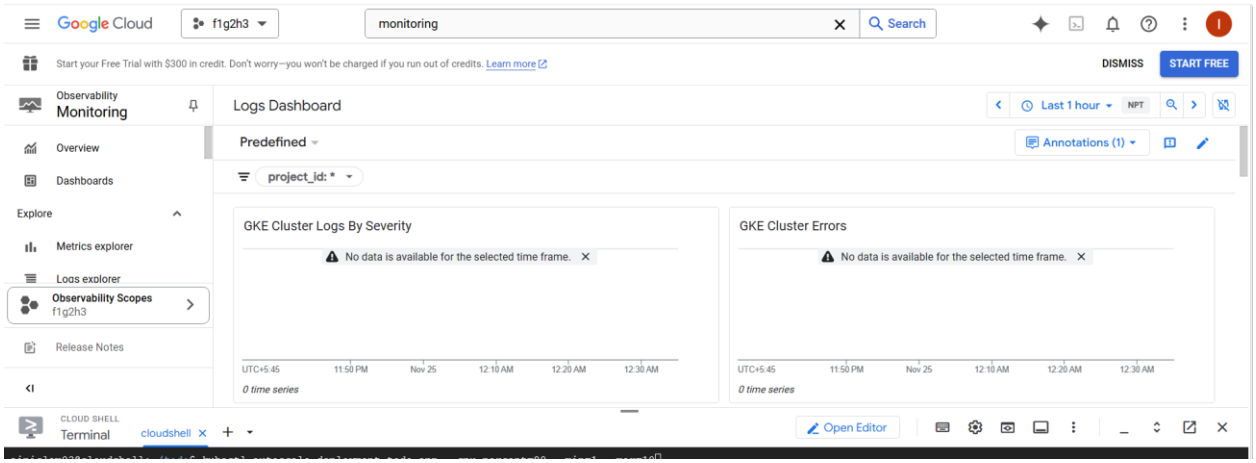

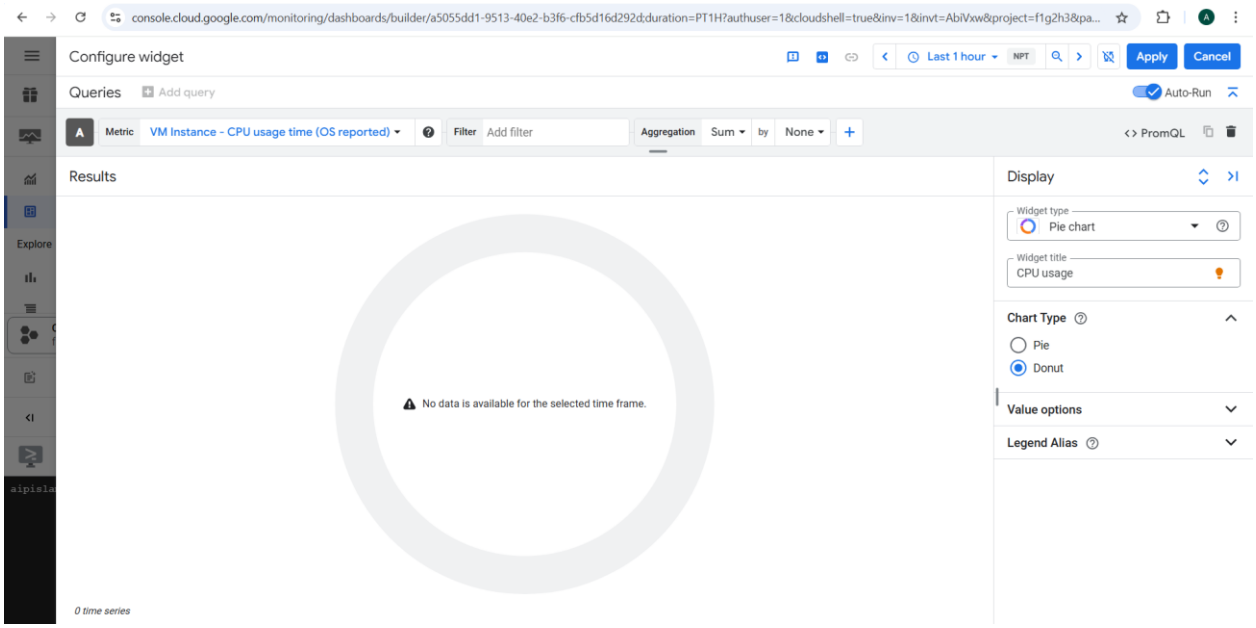

6. Cost Optimization

We can use Cloud Billing Reports to identify high-cost resources or underutilized resources. We can also use preemptive VMs and track before/after costs in Cloud Biling Reports

Conclusion

In this exercise, we explored essential practices for securing and scaling applications on Google Cloud, focusing on robust design, automation, and incident preparedness. For security, we implemented measures like IAM policies, encryption, and audit logs, complemented by a detailed Incident Response Plan. For scalability, we designed applications leveraging serverless and containerized approaches, implemented load balancing, auto-scaling, and monitored performance to optimize cost-effectiveness.

Recommendations

We can implement regular security testing and optimize our auto-scaling policies